



Timbre


Clean audio for transcription, straight from your Mac



Stop pointing your microphone at your own speakers. Timbre records what is playing inside your computer, whether that is a YouTube livestream, a Zoom call or a hearing, and gives you a clean MP3.

Everything runs on your Mac. Nothing is uploaded anywhere.

SCMP newsroom only · please do not share



A STEP-BY-STEP GUIDE · NO CODING REQUIRED

Why I built this

I have been working on this project for a few months. Like most of you, I often need a recording of something playing on my screen, whether it is a YouTube livestream, a Zoom or Google Meet call, or a hearing being streamed, so I can run it through Otter and get a transcript.

The only way I knew was to play the audio out loud and let my Mac's microphone pick it up. It worked, but badly.

What I was doing before

- My speakers had to stay on and loud, which bothered whoever was sitting near me
- The room went into the recording as well, including conversations and typing
- A phone call mirrored to my Mac would cut right into the audio
- Quality was never good, because I was recording a recording

What Timbre does instead

- Takes the sound from the sound card itself, before it reaches any speaker
- Nothing from the room gets in, only the audio you are capturing
- Volume stops mattering. Keep it low, use headphones, or mute it entirely
- Same quality as the original source, ready to upload for transcription

So I set out to build something better. Working with Claude Code and other AI tools to debug Swift, I spent months refining it until I had a Mac app that runs properly and does exactly what I needed.

Once I reached a version I was happy with, I cleaned it up, translated the whole interface into English, since my own version was in Brazilian Portuguese, and wrote this guide so the rest of the bureau can use it too.

Everything happens on your computer. No internet connection, no cloud service, no account, no upload. The recording is a file on your Desktop and nobody else ever sees it, which matters when the material is sensitive or embargoed.

It works with any source and any browser, including Chrome, Safari, Firefox, desktop apps and media players. If your Mac can play it, Timbre can record it.

How you will use it

Click the icon before the talk starts. Listen to it normally while it records. Click again when it ends. An MP3 appears on your Desktop, ready to use.

You need macOS 14.4 or later. Setup takes about five minutes and you only have to do it once.

1 Open Terminal

Press **⌘ + Space**, type **Terminal** and press **Return**.



A window with plain text will open. You only need it for the setup. After this you will never have to open it again.

2 Copy and paste these two lines

Copy the first line, paste it into Terminal and press **Return**. Wait until it finishes, then do the same with the second one.

```
git clone https://github.com/ipatrickbr/timbre.git ~/timbre
```

```
cd ~/timbre && ./vendor/build-lame.sh && ./build.sh
```

The second line takes a minute or two. A lot of text will scroll past, which is normal. It is done when you see the word **installed** followed by a path.

If macOS asks you to install the **command line developer tools**, click **Install**, wait for it to finish and run the line again.

3 Give Timbre permission

Open **Timbre** from Launchpad. The first time you run it, macOS will ask for permission to record system audio. Go to **System Settings**, then **Privacy & Security**, then **Screen & System Audio Recording**, and switch **Timbre** on.



Switch on the toggle circled in red. Timbre appears under **System Audio Recording Only**, meaning it can hear your Mac's audio and nothing else. It cannot see your screen.

4 Find the icon


Timbre sits in the menu bar at the top of your screen, as a small waveform.



The one circled in red is Timbre. Click it to start recording.

5 Choose how to record

Clicking the icon asks you one question.



Record while listening is what you will want almost always. You hear the talk normally and it records at the same time.


Record silently mutes your speakers while recording. Useful when you need the audio but not the noise.

6 While it records

The icon turns into a recording dot with a running clock.



Click it again at any time to pause or to finish.




Pausing leaves that stretch out of the file and picks up again in the same recording. Useful when a hearing breaks for recess.

If the icon says **no sound** instead of a clock, nothing is playing on your Mac yet. Press play on the stream and the clock will start.

7 Wait a moment for the file


When you finish, Timbre turns the recording into an MP3 and shows how long it will take.



A two-hour hearing takes under a minute. The folder opens by itself when it is ready.

8 Where your recordings are

Everything goes into a folder called **Timbre** on your Desktop. It opens automatically at the end of every recording, with the new file already selected.



Files are named by date and time. Double-click to play, or drag into your transcription tool.

9 Automatic transcription with Whisper

Timbre can also transcribe a recording the moment it finishes, without sending anything anywhere. This part is optional and you install it separately.

What Whisper is. Whisper is a speech recognition model built by OpenAI and released as open source, which means anyone can run it freely on their own machine. It is the same family of technology behind most transcription services, except here it runs on your laptop instead of somebody's server. Timbre uses `whisper.cpp`, a version rewritten to take advantage of the graphics chip in Apple Silicon Macs.

Languages. Whisper handles just under a hundred languages, including English, Portuguese, Spanish, French, German, Mandarin, Cantonese, Japanese, Korean, Arabic and Russian. You do not have to tell it which one you are recording, as it works that out on its own. It can also translate what it hears into English.

Why this matters for us. Sending a recording to Otter, DeepSeek or any similar service means handing our raw material to a third party. Some of what we record is sensitive, embargoed or comes from sources who would not appreciate it sitting on someone else's servers. Whisper removes that problem, because the audio never leaves the computer. It also removes two practical annoyances. There are no credits to run out of in the middle of a deadline, and it costs the newsroom nothing per hour of audio.


10 Turning transcription on

Run this line once in Terminal, the same way you did the setup.

```
cd ~/timbre && ./vendor/build-whisper.sh
```


It downloads the language models, about 2 GB in total, so give it a few minutes on a decent connection. You only do this once.

From then on, whenever a recording finishes, Timbre asks whether you want a transcript and tells you roughly how long it will take.



Say no and you just get the MP3, exactly as before.

Say yes and a progress bar shows how far along it is and roughly how much longer it will take. You can carry on working while it runs.




When it finishes, the folder opens with the audio and the transcript side by side, both in the same Timbre folder as the recording.

You get two files. The **.txt** is the plain transcript, for reading or pasting into a story. The **.srt** is the subtitle format used by video and audio editors such as Premiere, Final Cut and DaVinci Resolve. Drop it in alongside your footage and the words appear as timed captions, which saves a lot of work when you are cutting clips or publishing video with subtitles.

11 Recordings in another language

Whisper works out which language it is hearing on its own. If the recording turns out not to be in English, Timbre offers to produce an English version as well.



You end up with both: the transcript in the original language, and a second one marked **(English)** in the filename.

Use translations to navigate, not to quote. Machine translation is reliable enough to tell you what was said and to find the passage you need. It is not reliable for direct quotation. In testing, a Portuguese witness statement came back with "testimony" rendered as "appointment" and "quarter" as "semester". Go back to the original audio for anything that will appear between quotation marks.

12 What your Mac needs for this

Processor	Apple Silicon (M1 or newer). Intel Macs will run it, but far more slowly.
Memory	16 GB recommended. 8 GB works with a smaller model.
Disk space	About 2.5 GB for the engine and the two language models, one for transcription and a smaller one for translation.
Internet	Only for the one-time download. Transcription itself is offline.

How long it takes. On an Apple Silicon Mac with 16 GB of memory, transcription runs roughly eight times faster than the audio itself. One hour of a hearing comes back in about seven minutes. You can keep working while it runs.

This part is still in beta. Transcription is the newest piece of Timbre and I am still learning where the model struggles, whether that is heavy accents, crosstalk, poor stream audio or technical vocabulary. It will improve as we test it on real material. Please tell me what you run into, on Slack or by email. My details are on the last page.

? If something goes wrong

The recording came out silent. Permission is almost always the cause. Go back to step 3, because macOS gives you silence rather than an error when the permission is missing.

No file appeared. Nothing was playing during the recording. Timbre does not create a file when there is no sound at all.

The icon vanished after restarting. Open Timbre from Launchpad again. To have it there every time, add it under **System Settings, General, Login Items**.

The file is shorter than the time I recorded. That is expected. Recording starts at the first sound, so silence before the talk begins is not included. Nothing is lost.

Please keep this inside the newsroom. I built Timbre specifically for SCMP colleagues and it is not a public product. Please do not reproduce, copy or pass the app on to anyone outside the newsroom.

Questions, problems or ideas? Message me on Slack, or email igor.patrick@scmp.com. If something does not work, tell me which step you were on, as that is usually enough for me to fix it.