Timbre

Clean audio for transcription, straight from your Mac

Stop pointing your microphone at your own speakers. Timbre records what is playing inside your computer, whether that is a YouTube livestream, a Zoom call or a hearing, and gives you a clean MP3.

Everything runs on your Mac. Nothing is uploaded anywhere.

SCMP newsroom only · please do not share

A STEP-BY-STEP GUIDE · NO CODING REQUIRED

? Why I built this

I have been working on this project for a few months. Like most of you, I often need a recording of something playing on my screen, whether it is a YouTube livestream, a Zoom or Google Meet call, or a hearing being streamed, so I can run it through Otter and get a transcript.

The only way I knew was to play the audio out loud and let my Mac's microphone pick it up. It worked, but badly.

What I was doing before

- My speakers had to stay on and loud, which bothered whoever was sitting near me
- The room went into the recording as well, including conversations and typing
- A phone call mirrored to my Mac would cut right into the audio
- Quality was never good, because I was recording a recording

What Timbre does instead

- Takes the sound from the sound card itself, before it reaches any speaker
- Nothing from the room gets in, only the audio you are capturing
- Volume stops mattering. Keep it low, use headphones, or mute it entirely
- Same quality as the original source, ready to upload for transcription

So I set out to build something better. Working with Claude Code and other AI tools to debug Swift, I spent months refining it until I had a Mac app that runs properly and does exactly what I needed.

Once I reached a version I was happy with, I cleaned it up, translated the whole interface into English, since my own version was in Brazilian Portuguese, and wrote this guide so the rest of the bureau can use it too.

> **Everything happens on your computer.** No internet connection, no cloud service, no account, no upload. The recording is a file on your Desktop and nobody else ever sees it, which matters when the material is sensitive or embargoed.

It works with any source and any browser, including Chrome, Safari, Firefox, desktop apps and media players. If your Mac can play it, Timbre can record it.

? How you will use it

Click the icon before the talk starts. Listen to it normally while it records. Click again when it ends. An MP3 appears on your Desktop, ready to use.

> You need macOS 14.4 or later. Setup takes about five minutes and you only have to do it once.

1 Install it

Open **Timbre.dmg**. A window appears with the app and a shortcut to your Applications folder. Drag **Timbre** onto **Applications**, and that is the installation done.

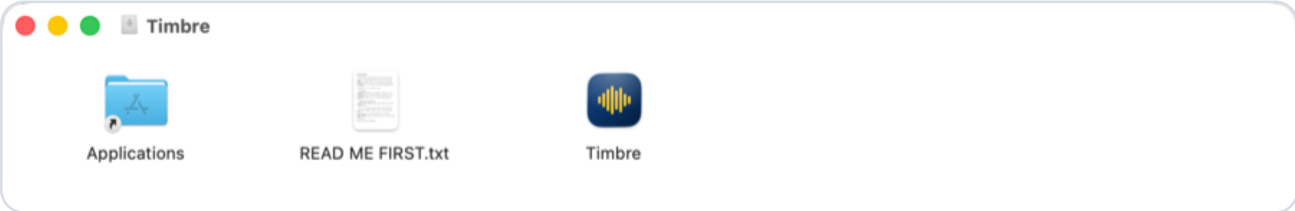

No terminal, no compiling, nothing to configure. The whole app is 17 MB.

> **Do not double-click Timbre in that window.** Opening it from the disk image cannot work, because the image is read-only and macOS cannot record permission to run something it cannot write to. You get a refusal with no way to allow it, which looks like the app is broken. Drag it to Applications first, then open it from Launchpad.

2 Getting macOS to open it

The first time you open Timbre, macOS refuses:

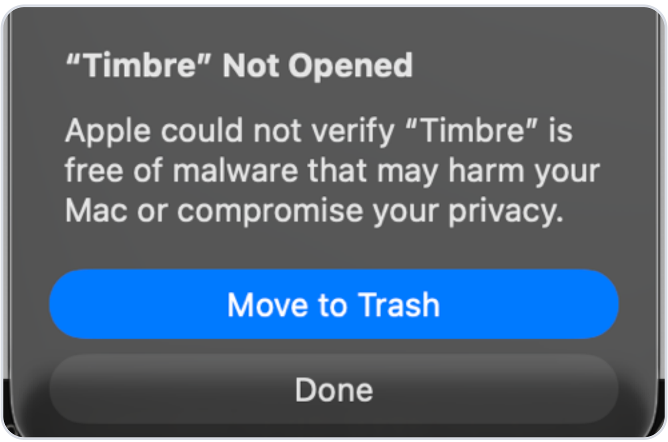

> **Click "Done", never "Move to Trash".** The blue button deletes the app, and it is the one that responds if you press Return without reading.

This happens because Timbre is not signed with a paid Apple developer certificate. On recent versions of macOS there is often no button anywhere to allow it, so the reliable way through is one line in Terminal.

Press **⌘ + Space**, type **Terminal**, press **Return**. Then paste this and press **Return** again:

```
xattr -dr com.apple.quarantine /Applications/Timbre.app
```

Nothing appears to happen. That is what success looks like. Open Timbre from Launchpad and it will start normally. You only do this once.

Some Macs do show a button instead, under **System Settings** › **Privacy & Security**, at the bottom. If you see it, clicking it works too:

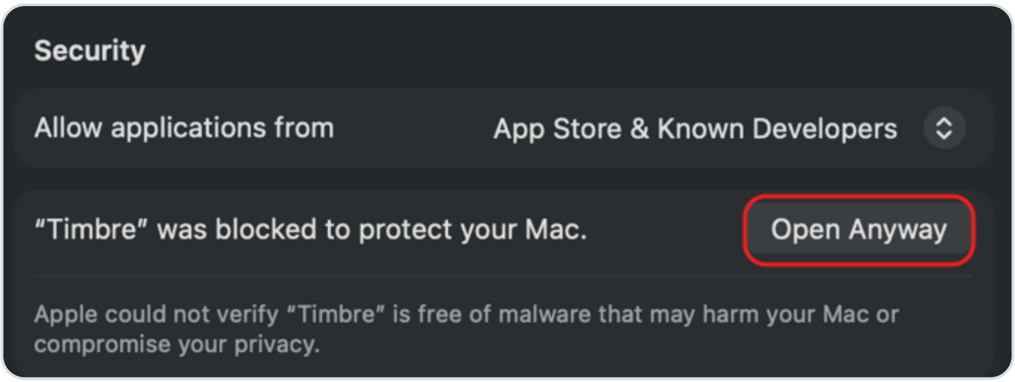

Either route gets you to the same place, and neither changes anything else on your Mac.

3 Give Timbre permission

Open **Timbre** from Launchpad. The first time you run it, macOS will ask for permission to record system audio. Go to **System Settings**, then **Privacy & Security**, then **Screen & System Audio Recording**, and switch **Timbre** on.

Switch on the toggle circled in red. Timbre appears under **System Audio Recording Only**, meaning it can hear your Mac's audio and nothing else. It cannot see your screen.

4 Find the icon

Timbre sits in the menu bar at the top of your screen, as a small waveform.

The one circled in red is Timbre. Click it to start recording.

5 Choose how to record

Clicking the icon asks you one question.

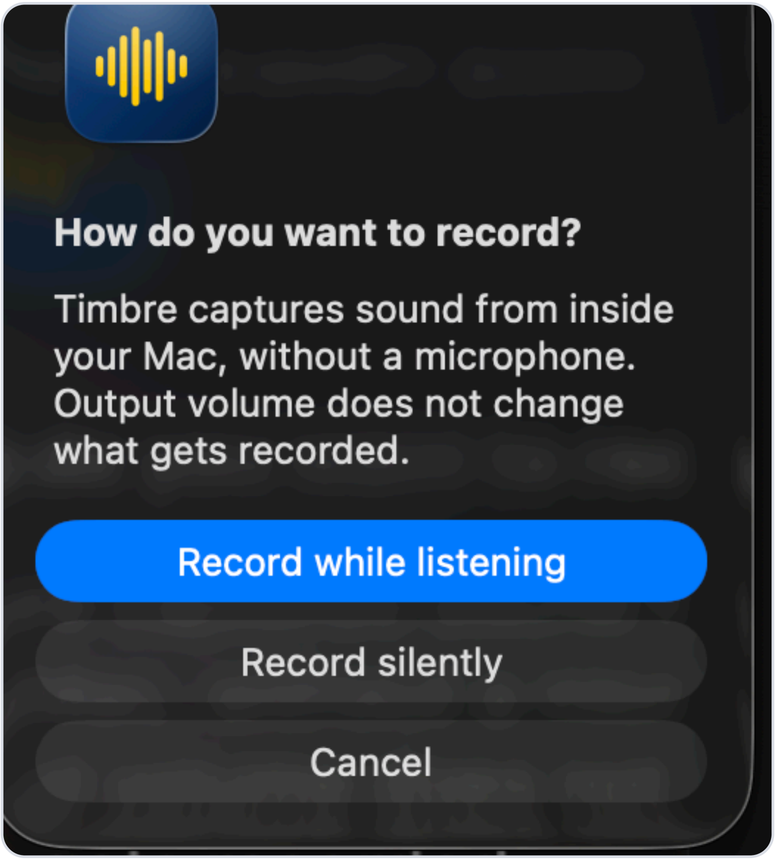

Record while listening is what you will want almost always. You hear the talk normally and it records at the same time.

Record silently mutes your speakers while recording. Useful when you need the audio but not the noise.

6 While it records

The icon turns into a recording dot with a running clock.

Click it again at any time to pause or to finish.

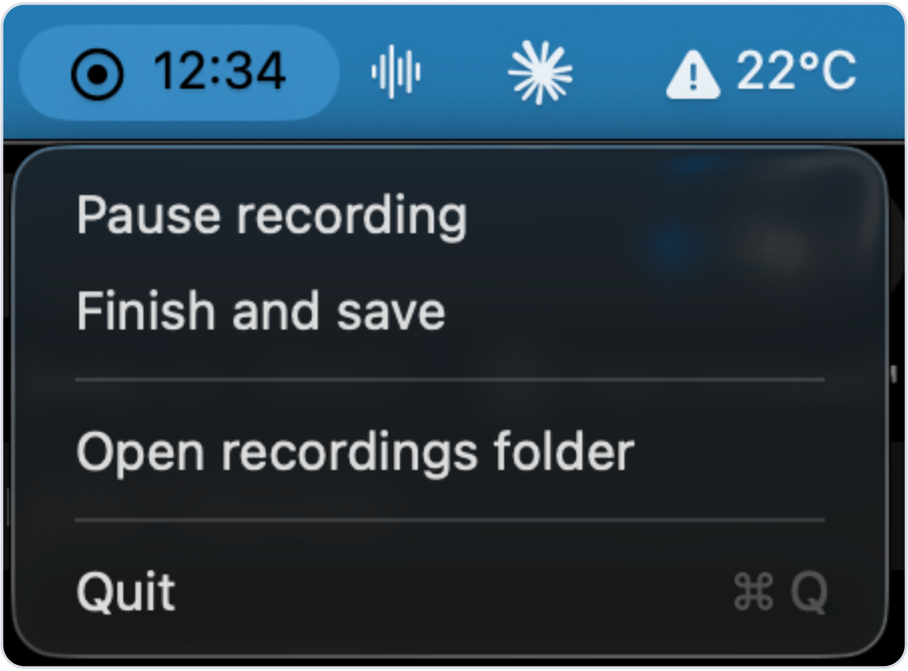

Pausing leaves that stretch out of the file and picks up again in the same recording. Useful when a hearing breaks for recess.

> If the icon says **no sound** instead of a clock, nothing is playing on your Mac yet. Press play on the stream and the clock will start.

7 Wait a moment for the file

When you finish, Timbre turns the recording into an MP3 and shows how long it will take.

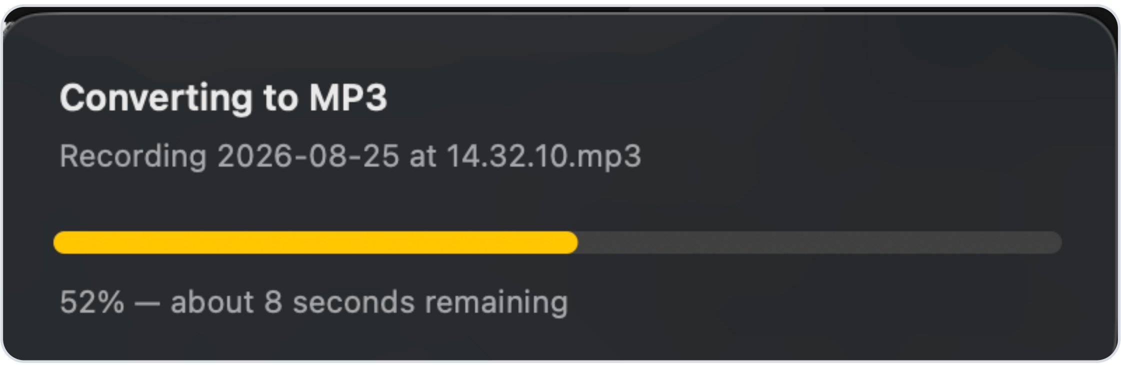

A two-hour hearing takes under a minute. The folder opens by itself when it is ready.

8 Where your recordings are

Everything goes into a folder called **Timbre** on your Desktop. It opens automatically at the end of every recording, with the new file already selected.

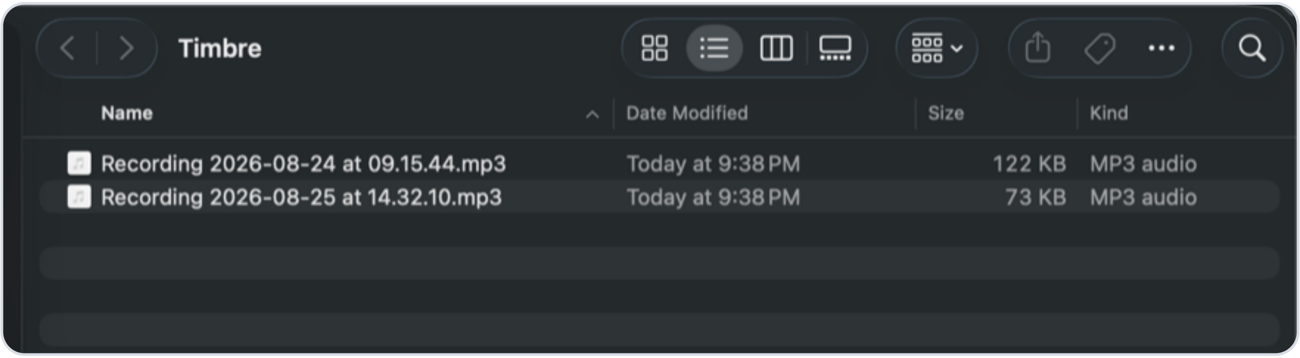

Files are named by date and time. Double-click to play, or drag into your transcription tool.

9 Automatic transcription with Whisper

Timbre can also transcribe a recording the moment it finishes, without sending anything anywhere. This part is optional and you install it separately.

What Whisper is. Whisper is a speech recognition model built by OpenAI and released as open source, which means anyone can run it freely on their own machine. It is the same family of technology behind most transcription services, except here it runs on your laptop instead of somebody's server. Timbre uses whisper.cpp, a version rewritten to take advantage of the graphics chip in Apple Silicon Macs.

Languages. Whisper handles just under a hundred languages, including English, Portuguese, Spanish, French, German, Mandarin, Cantonese, Japanese, Korean, Arabic and Russian. You do not have to tell it which one you are recording, as it works that out on its own. It can also translate what it hears into English.

Why this matters for us. Sending a recording to Otter, DeepSeek or any similar service means handing our raw material to a third party. Some of what we record is sensitive, embargoed or comes from sources who would not appreciate it sitting on someone else's servers. Whisper removes that problem, because the audio never leaves the computer. It also removes two practical annoyances. There are no credits to run out of in the middle of a deadline, and it costs the newsroom nothing per hour of audio.

10 Turning transcription on

Nothing to install. The first time you say yes to a transcript, Timbre asks whether it can fetch the language models it needs.

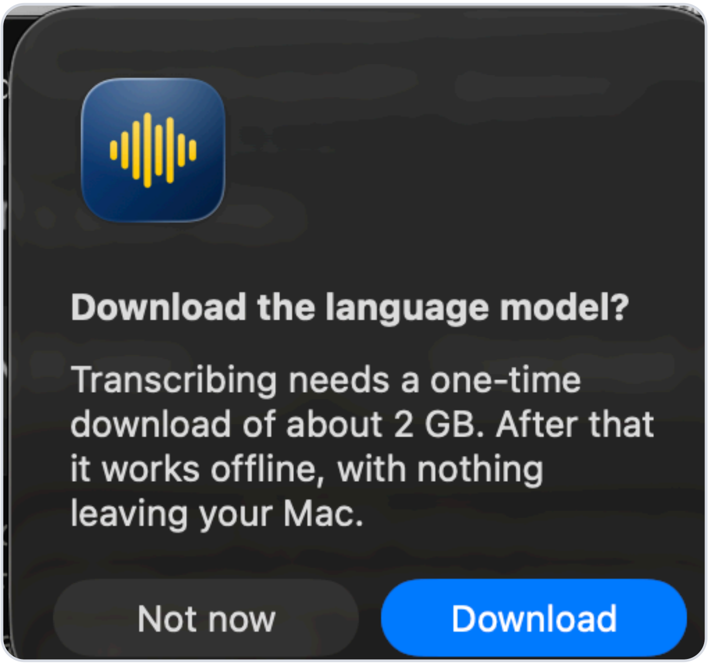

Roughly 2 GB, and it happens once. After this, transcription never touches the internet again.

Say yes and it downloads them itself, showing how far along it is.

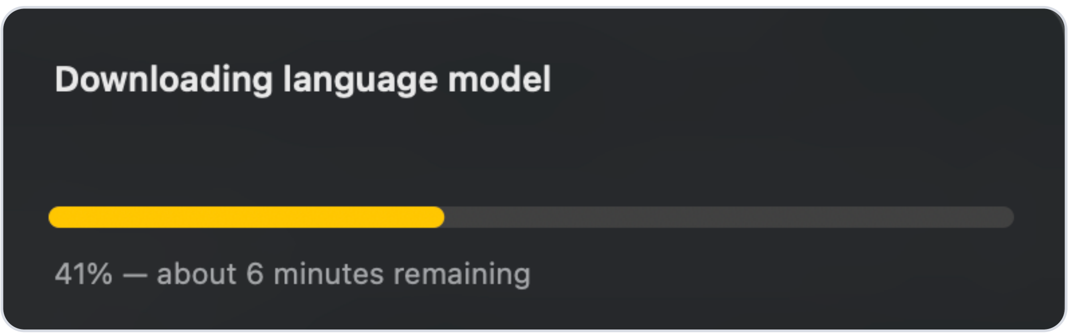

Give it a few minutes on a decent connection. You can keep working, and the transcript starts as soon as the download finishes.

From then on, whenever a recording finishes, Timbre asks whether you want a transcript and tells you roughly how long it will take.

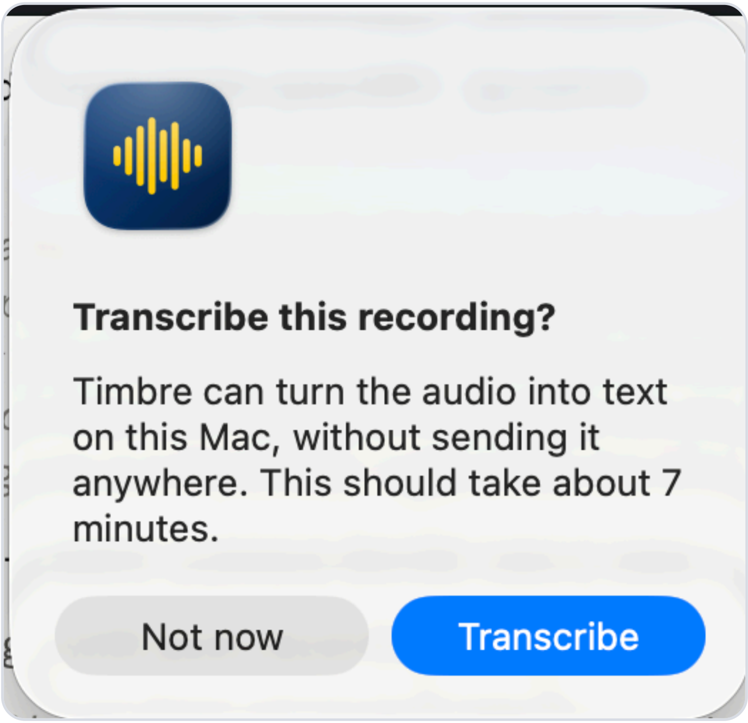

Say no and you just get the MP3, exactly as before.

Say yes and a progress bar shows how far along it is and roughly how much longer it will take. You can carry on working while it runs.

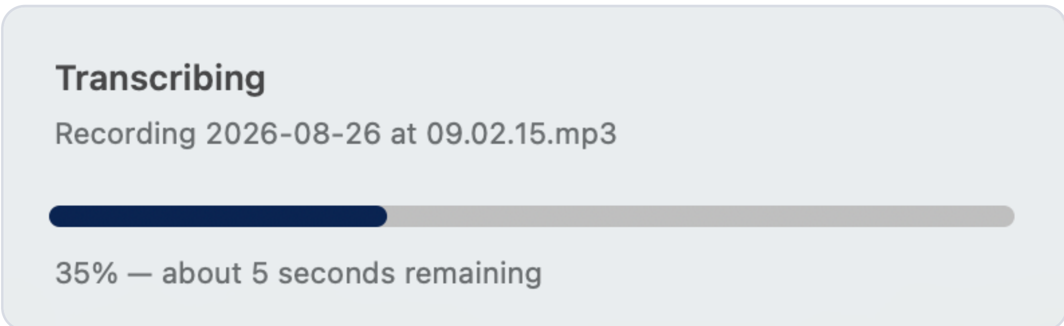

When it finishes, the folder opens with the audio and the transcript side by side, both in the same Timbre folder as the recording.

You get two files. The **.txt** is the plain transcript, for reading or pasting into a story. The **.srt** is the subtitle format used by video and audio editors such as Premiere, Final Cut and DaVinci Resolve. Drop it in alongside your footage and the words appear as timed captions, which saves a lot of work when you are cutting clips or publishing video with subtitles.

11 Recordings in another language

Whisper works out which language it is hearing on its own. If the recording turns out not to be in English, Timbre offers to produce an English version as well.

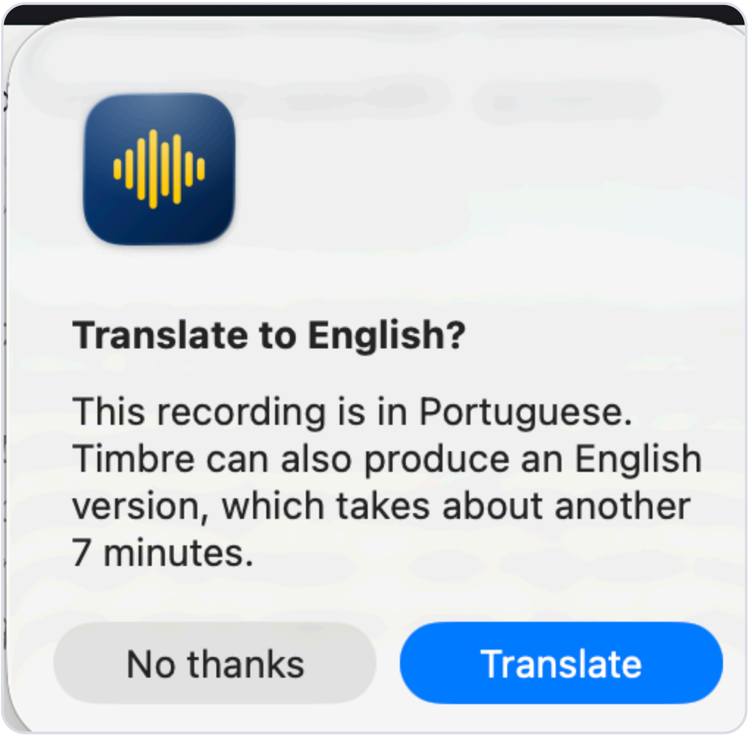

You end up with both: the transcript in the original language, and a second one marked **(English)** in the filename.

> **Use translations to navigate, not to quote.** Machine translation is reliable enough to tell you what was said and to find the passage you need. It is not reliable for direct quotation. In testing, a Portuguese witness statement came back with "testimony" rendered as "appointment" and "quarter" as "semester". Go back to the original audio for anything that will appear between quotation marks.

12 Using this within our editorial policy

Timbre uses AI in two places, the transcript and the translation, so the **SCMP Editorial Generative AI Policy** applies to both. A few things are worth stating plainly.

Where Timbre already works in our favour.

What the policy asks for

- Tools that protect our material best, and corporate accounts rather than free ones that may be read by moderators or used for training
- A proper reference for anything an AI tool surfaced
- A human check before AI output reaches readers

How Timbre answers it

- Nothing is uploaded at all, so there is no third party to trust in the first place and no account to worry about
- The audio, the transcript and the timestamps stay together, so the original is always one click away
- The **.srt** points to the exact second, turning that check into a few seconds of work instead of a hunt through an hour of tape

It also improves the raw material. Recording from the sound card instead of a microphone removes room noise, phone interruptions and the quality loss of recording a recording, so there is less for the transcription to get wrong. And because it costs nothing per hour, there are no credits to run out of on deadline.

What still falls to you.

The transcript is not a source. Under the policy, nothing produced by an AI tool can be cited as a primary source. The recording is the source. The transcript is a way of finding your way around it. Timbre deliberately keeps the audio next to the text so the original is always within reach.

Check before you quote. Nothing from an AI tool should reach readers without a human checking it first. Before a line goes into a story between quotation marks, listen to that passage in the MP3 and confirm it word for word. The **.srt** file gives you the timestamp, so finding the moment takes seconds.

Tell your editor. The policy asks you to inform your reporting manager when AI has been used in your work, and says nobody in the publishing chain should be caught by surprise. If a Timbre transcript fed into a story, mention it when you file.

On data protection. The policy asks us to prefer tools that protect our material best. Timbre runs entirely on your own machine. No audio is uploaded, no account is involved, and nothing you record is seen by a third party or used to train a model. That is the main reason it was built this way instead of leaning on a cloud service.

> **Timbre is in testing.** The SCMP Editorial Generative AI Policy requires tools outside the approved whitelist to be cleared by a senior manager, and that clearance is not yet in place. Check with your editor before using Timbre on anything headed for publication.

13 What your Mac needs for this

Processor	Apple Silicon (M1 or newer). Intel Macs will run it, but far more slowly.
Memory	16 GB recommended. 8 GB works with a smaller model.
Disk space	About 2.5 GB for the engine and the two language models, one for transcription and a smaller one for translation.
Internet	Only for the one-time download. Transcription itself is offline.

How long it takes. On an Apple Silicon Mac with 16 GB of memory, transcription runs roughly eight times faster than the audio itself. One hour of a hearing comes back in about seven minutes. You can keep working while it runs.

This part is still in beta. Transcription is the newest piece of Timbre and I am still learning where the model struggles, whether that is heavy accents, crosstalk, poor stream audio or technical vocabulary. It will improve as we test it on real material. Please tell me what you run into, on Slack or by email. My details are on the last page.

? If the disk image will not work

There is a longer way that builds Timbre from source on your own machine. It takes about five minutes and needs the Terminal, but it sidesteps the security warning entirely, because an app you compiled yourself is not treated as something downloaded from the internet.

Press **⌘ + Space**, type **Terminal** and press **Return**. Then copy these lines in one at a time, waiting for each to finish.

```
git clone https://github.com/ipatrickbr/timbre.git ~/timbre
```

```
cd ~/timbre && ./vendor/build-lame.sh && ./build.sh
```

A lot of text will scroll past, which is normal. It is done when you see the word **installed** followed by a path. Timbre then appears in Launchpad and you can carry on from step 3.

> If macOS offers to install the **command line developer tools**, click **Install**, wait for it to finish and run the line again.

Transcription still downloads its models by itself, exactly as described above. If you would rather build that part too, there is a third line for it:

```
cd ~/timbre && ./vendor/build-whisper.sh
```

? If something goes wrong

The recording came out silent. Permission is almost always the cause. Go back to step 3, because macOS gives you silence rather than an error when the permission is missing.

No file appeared. Nothing was playing during the recording. Timbre does not create a file when there is no sound at all.

The icon vanished after restarting. Open Timbre from Launchpad again. To have it there every time, add it under **System Settings**, **General**, **Login Items**.

The file is shorter than the time I recorded. That is expected. Recording starts at the first sound, so silence before the talk begins is not included. Nothing is lost.

Please keep this inside the newsroom. I built Timbre specifically for SCMP colleagues and it is not a public product. Please do not reproduce, copy or pass the app on to anyone outside the newsroom.

Questions, problems or ideas? Message me on Slack, or email **igor.patrick@scmp.com**. If something does not work, tell me which step you were on, as that is usually enough for me to fix it.